

Remote Procedure Call (RPC)

Web Applications and Services
Spring term

Naercio Magaia




Contents

- Distributed systems
- Local Procedure Call
- Remote Procedure Call
- Representing data
- Interface Definition Language
- RPC Semantics
- Google Protocol Buffers
- Apache Thrift

Distributed systems

- It is a group of computers working together as to appear as a single computer to the end-user
- These machines have a shared state, operate concurrently
- They communicate via a network to achieve a task
- Networked computers communicate and coordinate their actions only via messages passing



Why distributed systems?


- The desire to share resources
- For scalability, performance and high availability of applications
- Some examples include
 - web search
 - distributed information
 - network file systems
 - real-time process control
 - parallel computation

Benefits of distributed systems

- Computation speedup, better performance
 - Parallel computation on multiple servers
 - Solve bigger problems
- Redundancy, reliability, fault-tolerant
 - Several machines can provide the same services, so if one is unavailable, the work does not stop.
- Many applications are inherently distributed

Challenges

- The network may fail
- Processes may crash
- Writing a program to run on a single computer is comparatively easy?
- How can we coordinate the processes of these machines?




Programming

- The goal of making the programming of distributed systems look similar, if not identical, to conventional programming

Examples



The black hole image



Perseverance on Mars


Local Procedure Call

- Everything in the same address space
- Calls cannot fail (i.e., they can but programmers have full control)

```
ret = foo(64, "string", &myStruct);
```
- Push arguments, local variables and return address in the stack
- Call foo
- Pop everything out of the stack and put ret to a register
- Continue with the next call
- What about parameter passing?

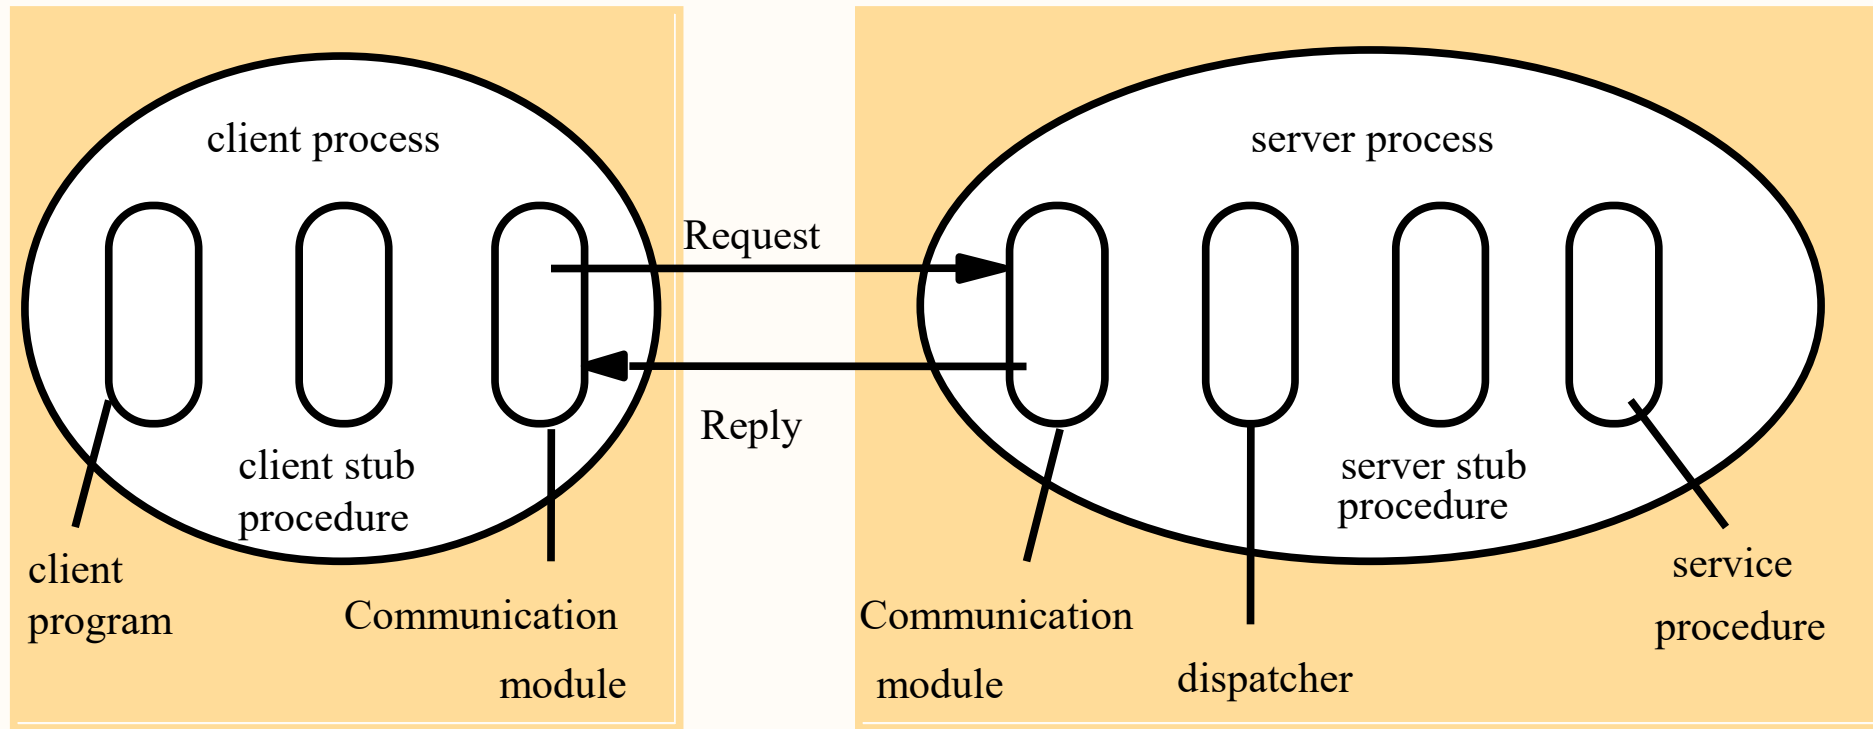
Remote Procedure Call (RPC)

- Makes a call to a remote function look the same as a local function call



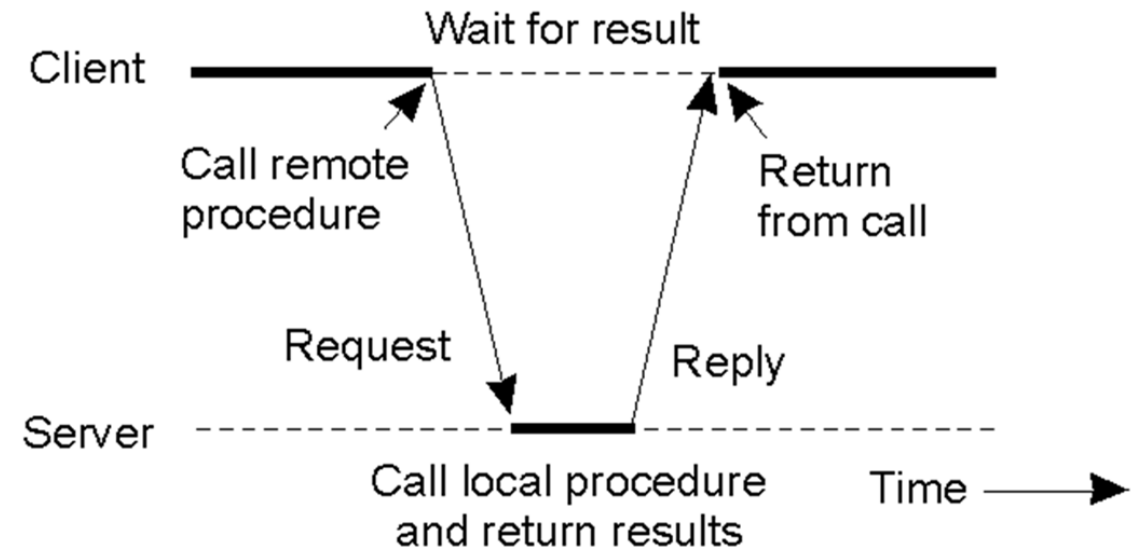
- In practice,
 - What if the service crashes during the function call?
 - What if a message is lost?
 - What if a message is delayed?
 - If something goes wrong, is it safe to retry?

RPC at a Glance



RPC

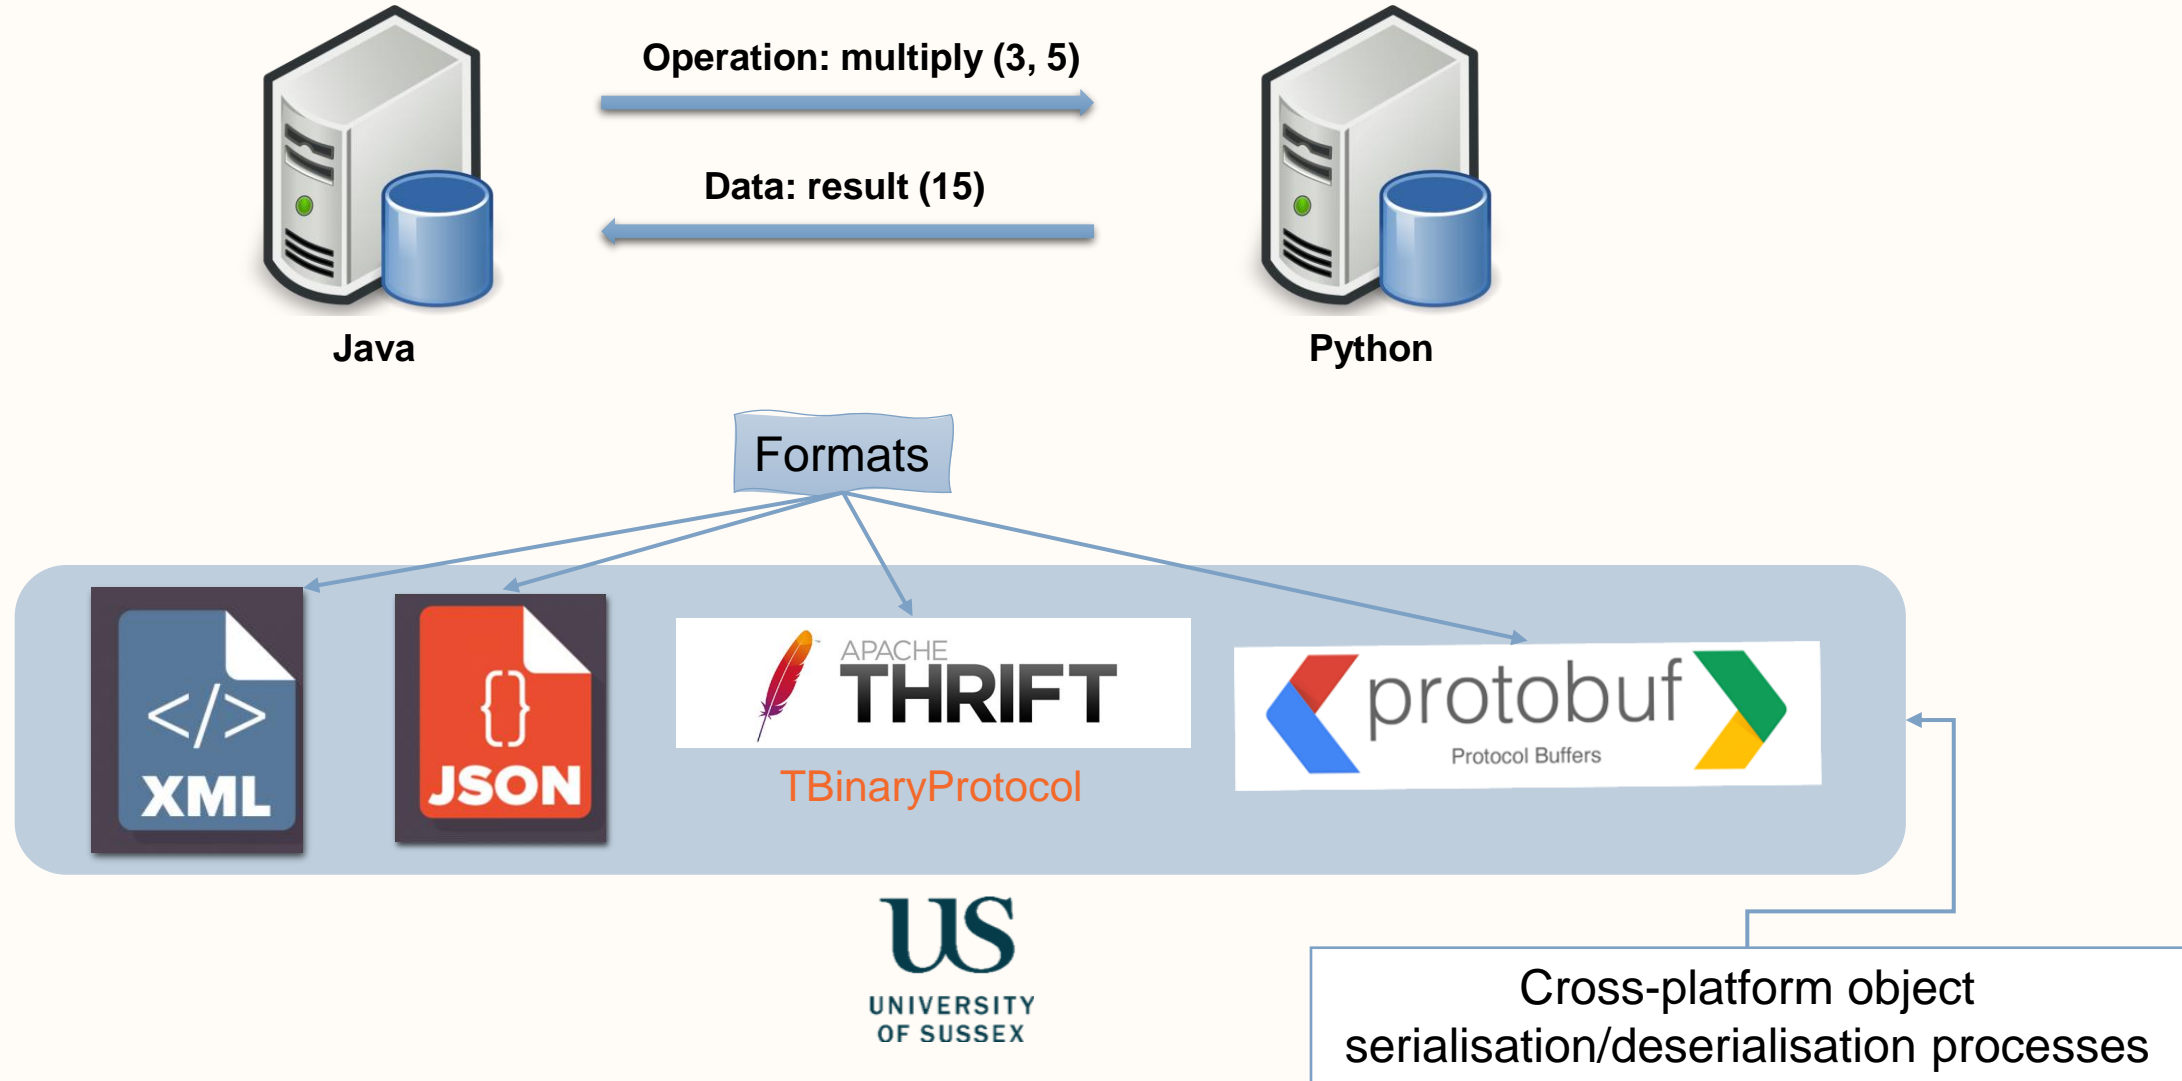
- A process on machine A can call a procedure on machine B
- The process on A is suspended and execution continues on B
- When B returns, the return value is passed to A, and it continues its execution.



RPC history


- SunRPC/ONC RPC (1980s, basis for NFS)
- CORBA: object-oriented middleware, hot in the 1990s
- Microsoft's DCOM and Java RMI (similar to CORBA)
- SOAP/XML-RPC: RPC using XML and HTTP (1998)
- Thrift (Facebook, 2007)
- gRPC (Google, 2015)
- REST (often with JSON)
- ...

RPC



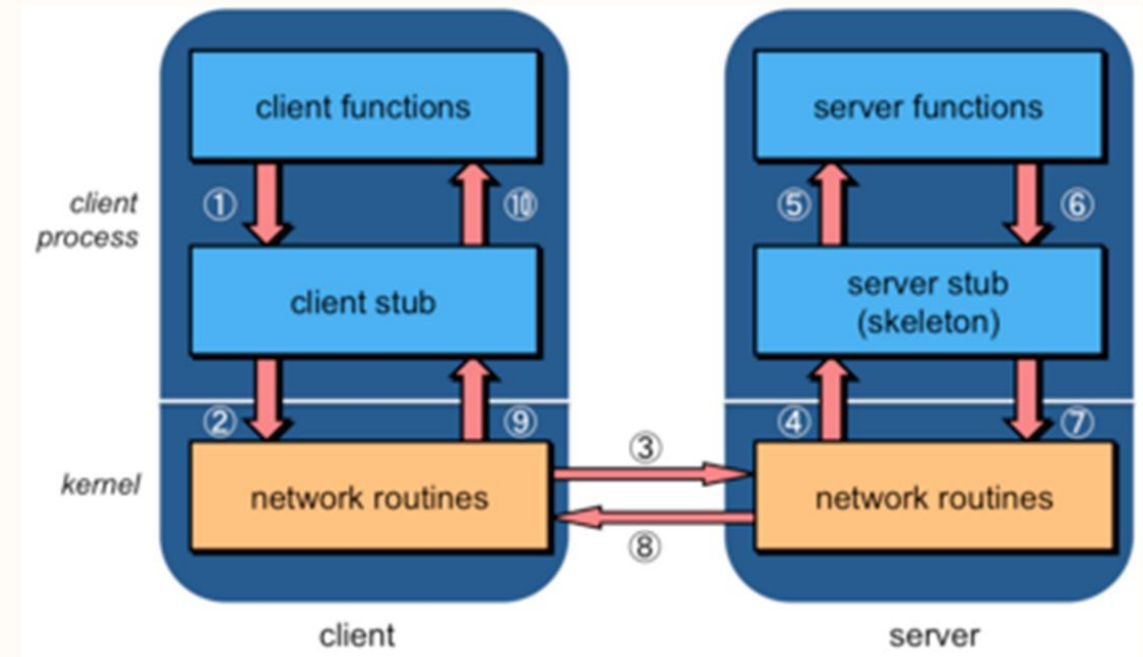
RPC : Client side

- Emulate the behaviour of local calls (transparently)
- Stub function
 - the same signature as if it was called locally
 - marshals parameters
 - sends request message
 - waits for a response from the server
 - un-marshals the response & returns the appropriate data



RPC : Server side

- Dispatcher
 - Receives client requests
 - Identifies appropriate function to call
- Skeleton
 - Un-marshals parameters
 - Calls the local function
 - Marshals the response & sends it back to the dispatcher



Parameter passing

- Passing by value (is simple)
 - Just copy the value into the network message
- Passing by reference (is hard)
 - It makes no sense to pass an address to a remote machine
 - A memory location in a process on one machine is meaningless to another process on another machine

Parameter passing

1. Copy items referenced to message buffer
2. Ship them over
3. Un-marshall data at server
4. Pass local pointer to server stub function
5. Send new values back

Representing data

- Local Calls
 - Primitive types have the same exact representation
- But, for Remote Calls
 - Different byte ordering
 - Big Endian – most significant byte is stored in the smallest address
 - Little Endian - least significant byte is stored in the smallest address
 - Different sizes of integers and other types
 - Different floating-point representations
 - Different Languages!!

Representing data

- We need standards!
 - SunRPC uses eXternal Data Representation (XDR)
 - Abstract Syntax Notation (ASN.1)
 - WSDL for Web Service
 - Google Protocol Buffers
 - JSON
- Implicit typing
 - only values are transmitted, not data types or parameter info
- Explicit typing
 - Type is transmitted with value

Interface Definition Language

- Allow programmer to specify remote procedure interfaces (i.e., names, parameters, return values)
- Pre-compiler can use this to generate
 - client and server stubs
 - Marshalling code
 - Un-marshalling code
 - Network transport routines

RPC Semantics

- Local Calls
 - exactly Once!
- But RPC involves networking and a remote host
 - delays, server failures
- A remote procedure call may be called
 - never: server crashed, or server process died before executing server code
 - once: everything worked well, as expected
 - more than once: Why?

RPC Semantics

- RPC systems may offer either
 - *maybe* semantics
 - *at least once* semantics
 - *at most once* semantics
- How can these be implemented?
- What are the ramifications with respect to
 - idempotent functions?
 - non-idempotent functions?

Protocol Buffers (Protobuf)

- Protobuf is Google's language-neutral, platform-neutral, extensible mechanism for serializing structured data for
 - communications protocols
 - data storage
 - specification of services
- Think XML, but smaller, faster, and simpler
- It currently support generated code in Java, Python, Objective-C, and C++
 - With the new proto3 language version, it is possible to also work with Kotlin, Dart, Go, Ruby, PHP, and C#

Protocol Buffers

- Why not XML?
 - Protobuf is simpler, i.e., 3 to 10 times smaller and 20 to 100 times faster
 - It generate data access classes that are easier to use programmatically (should we write source code and auto-produce description/IDL or vice versa?)

```
<person>
  <name>John Doe</name>
  <email>jdoes@example.com</email>
</person>
```

```
# Textual representation of a protocol buffer
# This is *not* the binary format used on the wire
person {
  name: "John Doe"
  email: "jdoes@example.com"
}
```

- Many projects use protocol buffers, including the following: [gRPC](#), [Google Cloud](#), [Envoy Proxy](#)

Protocol Buffers : RPC

- The protocol compiler will read .proto files and generate
 - an abstract interface called SearchService (must be implemented!)
 - a corresponding “stub” implementation
- The stub forwards all calls to an RpcChannel
- an abstract interface that must be defined for RPC system

```
message SearchRequest {  
    required string query = 1;  
    optional int32 page_number = 2;  
    optional int32 result_per_page = 3;  
}
```

```
message SearchResponse {  
    ...  
}
```

```
service SearchService {  
    rpc Search(SearchRequest) returns (SearchResponse);  
}
```

Apache Thrift

- A software framework for scalable cross-language services development
- Combines a software stack with a code generation engine
- Builds services that work efficiently and seamlessly between C++, Java, Python, PHP, Ruby, Erlang, Perl, Haskell, C#, Cocoa, JavaScript, Node.js, Smalltalk, OCaml and Delphi and other languages
- Type system consists of pre-defined base types, user-defined structs, container types, exceptions and service definitions

Thrift Types

- Base Types

- bool: A boolean value (true or false)
- byte: An 8-bit signed integer
- i16: A 16-bit signed integer
- i32: A 32-bit signed integer
- i64: A 64-bit signed integer
- double: A 64-bit floating point number
- string: A text string encoded using UTF-8 encoding

- Special Types

- binary: a sequence of unencoded bytes

- Containers

- list: An ordered list of elements
- set: An unordered set of unique elements
- map<type1,type2>: A map of strictly unique keys to values

- Container elements may be of any valid Thrift Type.

Thrift Types

- **Structs**

- define a common object – they are essentially equivalent to classes in object-oriented programming (OOP) languages, but without inheritance
- A struct has a set of strongly typed fields, each with a unique name identifier
- Fields may have various annotations (numeric field IDs, optional default values, etc.) that are described in the Thrift IDL

- **Exceptions**

- are functionally equivalent to structs, except that they inherit from the native exception base class

Thrift Types

- **Services**
 - Are defined using Thrift types
 - Definition of a service is semantically equivalent to defining an interface (or a pure virtual abstract class) in OOP
 - The Thrift compiler generates fully functional client and server stubs that implement the interface
 - A service consists of a set of named functions, each with a list of parameters and a return type

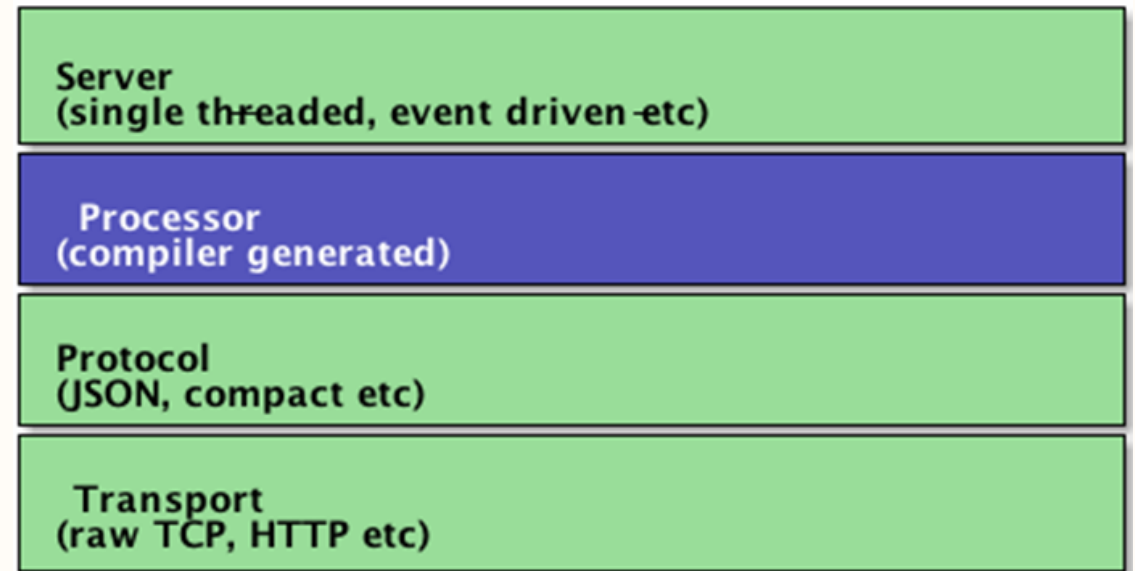
Thrift network stack

- **Transport**

- Provides a simple abstraction for reading/writing from/to the network
- Enables Thrift to decouple the underlying transport from the rest of the system (serialization/deserialization, for instance)

- **Protocol**

- Defines a mechanism to map in-memory data structures to a wire-format
 - specifies how datatypes use the underlying Transport to encode/decode themselves
- Some examples of protocols in this sense include JSON, XML, plain text, compact binary, etc.



Network Stack

- **Processor**

- Encapsulates the ability to read data from input streams and write to output streams
- The input and output streams are represented by Protocol objects

- **Server**

- Create a transport
- Create input/output protocols for the transport
- Create a processor based on the input/output protocols
- Wait for incoming connections and hand them off to the processor



Thrift definition file

```
service Calculator extends shared.SharedService {
```

```
/**
```

```
* A method definition looks like C code. It has a return type, arguments,  
* and optionally a list of exceptions that it may throw. Note that argument  
* lists and exception lists are specified using the exact same syntax as  
* field lists in struct or exception definitions.
```

```
*/
```

```
void ping(),
```

```
i32 add(1:i32 num1, 2:i32 num2),
```

```
i32 calculate(1:i32 logid, 2:Work w) throws (1:InvalidOperation ouch),
```

```
/**
```

```
* This method has a oneway modifier. That means the client only makes  
* a request and does not listen for any response at all. Oneway methods  
* must be void.
```

```
*/
```

```
oneway void zip()
```

```
}
```

Method definitions can be terminated using comma or semi-colon

void is a valid return type for functions

Arguments can be primitive types or structs, likewise for return types

Next Lecture ...

- ✓ Introduction
- ✓ HTTP, Caching, and CDNs
- ✓ Views
- ✓ Templates
- ✓ Forms
- ✓ Models
- ✓ Security
- ✓ Transactions
- ✓ Remote Procedure Call
- **Web Services**
 - RESTful Services
 - Time
 - Elections/Group Communication
 - Zookeeper