

# Elections and Group Communication

Web Applications and Services  
Spring Term

Naercio Magaia



# Contents

- Assumptions Failure Models
- Elections
- Group Communication

# Assumptions

- Processes are connected by reliable (not necessarily FIFO - TCP) channels among each other, although the network may be unreliable
  - In synchronous systems, there is an upper bound for the message transmission, processing delay, and clock drift
- Processes fail independently
  - A process failure is not a threat to other processes or communication among them
- Network can be partitioned

# Assumptions

- Reliability assumption
  - **eventually** any failed link or router will be repaired or circumvented
  - but processes may not all be able to communicate at the same time
- Processes only fail by crashing
  - Unless explicitly stated – arbitrary (or Byzantine) failures
- A correct process is the one that does not crash at any point during execution
- There is no correct process which then failed

# Failure Detector

- A service that processes queries about the state of a process
- Unreliable
  - may produce one of the *unsuspected* or *suspected* values when given the identity of a process (both hints)
- Reliable
  - may produce one of the *unsuspected* or *failed* values when given the identity of a process (*unsuspected* is a hint)
    - a crashed process stays crashed

# Failure Detector

- Implemented by sending heartbeats to the detector and timeouts
  - Trade-offs when selecting timeout values
  - What timeout value to select in a synchronous distributed system?

# Elections

- Is a helper algorithm to bootstrap other algorithms
- If the leader crashes or retires, a new election is started
- A process can start a single election but all processes in a system can start an election at the same time
- A process can be a participant (i.e., engaged in some election) or a non-participant
- The elected process has the 'largest' identifier

# Elections

- Each process has a variable  $electd_i$ , which can be undefined or contain the leader process
- Requirements
  - Safety: during an election a participant process'  $electd_i$  value is either undefined or  $P$ , where  $P$  is the elected leader process
  - Liveness: eventually all processes'  $electd_i$  value is defined
- Evaluation
  - Bandwidth: total number of sent messages to elect leader



# Ring-based Algorithm

- Processes arranged in a ring
  - communicate with next process in the ring
- Assumptions
  - no failures occur
  - The system is asynchronous
  - The elect process (i.e., coordinator) has the largest 'identifier'
- Initially, every process is marked as a non-participant in an election
- Any process can begin an election
  - marks itself as a participant

# Ring-based Algorithm

- Sends a message to the next process (with its identifier in it)
- When an **election message** is received, process identifiers are compared

**if**  $id_{message} > id_{process}$  **then**

forward message to the next process

**else if**  $id_{message} < id_{process}$  **then**

if receiver is not a participant,

substitute id and forward

**else**

discard message

# Ring-based Algorithm

- *On forwarding, mark itself as participant*
- If the received identifier is that of the receiver itself, then this process's identifier is the greatest
  - it becomes the leader
- The leader marks itself as a *non-participant* once more and sends an *elected message* to its neighbour, announcing its election and *id*

# Ring-based Algorithm

- When a process receives an elected message
  - it marks itself as a non-participant, sets its variable *elected<sub>i</sub>* to the *id* in the message
  - If not the coordinator, then forward elected message
- Note: duplicate messages from multiple elections are discarded as soon as possible
- Bandwidth: what's the worst and best case for a single election?

# The Bully Algorithm

- Processes can crash during election (fail - stop)
- The system is synchronous (i.e., timeouts to detect a process failure)
- Each process knows which processes have higher identifiers (IDs) and it can communicate with all these processes
- Election message
  - sent to announce an election

# The Bully Algorithm

- Answer message
  - sent in response to an election message
- Coordinator message
  - sent to announce the elected process
- Election begins when one or more processes identify a failed leader

# The Bully Algorithm

- The process with the highest ID can bully all other processes
- A process with a lower identifier can begin an election by sending an election message to processes with higher identifiers and waits for an answer
- If none arrives within time  $T$  (i.e., *synchronous system*), the process bullies the ones with lower IDs and announces itself as the leader

# The Bully Algorithm

- Otherwise, the process waits a further period for a *coordinator message* to arrive from the new coordinator
- If none arrives (i.e., a process crashed), it begins another election
- If a process receives a *coordinator message*, it sets its variable *elected<sub>i</sub>* to the ID of the coordinator
- If a process receives an election message, it sends back an answer message and begins another election
  - unless it has begun one already



# The Bully Algorithm


- Bandwidth

- Best case

- the process with the second-highest identifier notices the coordinator's failure ( $N - 2$  coordinator messages)

- Worst case

- $O(N^2)$  messages in the worst case
- the process with the lowest identifier detects the coordinator's failure ( $N - 1$  processes begin elections)



# Group Communication

- Collection of processes can communicate reliably over one-to-one channels and may fail only by crashing (fail-stop)
- Processes are members of groups – destinations of messages sent with the multicast operation – a process belongs to a single group
- ***multicast( $g, m$ )***: sends message  $m$  to all members of group  $g$
- ***deliver( $m$ )***: delivers message  $m$  sent by multicast to the process

# Group Communication

- A multicast message can be queued before the *deliver(m)* operation is called
- Every message *m* carries the unique identifier of the sender and the unique destination group identifier

# Basic Multicast

- Primitive multicast service
  - a correct process will eventually deliver the message as long as the sender does not crash
- Implemented by multi-unicasting the message using reliable channels
- *B-multicast*( $g, m$ ): *send*( $m$ ) to each process belonging to  $g$  separately
- On *receive*( $m$ ) at  $p$ : *B-deliver*( $m$ ) at  $p$
- Feedback implosion by acknowledgments when calling *send*( $m$ )

# Reliable Multicast

- Integrity
  - a correct process  $p$  delivers a message  $m$  at most once
- Validity
  - if a correct process multicasts message  $m$ , then it will eventually deliver  $m$
- Agreement
  - if a correct process delivers message  $m$ , then all other correct processes in  $g$  will eventually deliver  $m$  (all-or-nothing)

# Reliable Multicast

*On initialization*

*Received* := {};

*For process p to R-multicast message m to group g*

*B-multicast(g, m);*      //  $p \in g$  is included as a destination

*On B-deliver(m) at process q with g = group(m)*

*if* ( $m \notin \text{Received}$ )

*then*

*Received* := *Received*  $\cup$  { *m* } ;

*if* ( $q \neq p$ ) *then B-multicast(g, m); end if*

*R-deliver m;*

*end if*

- Inefficient algorithm
- each message is sent  $|g|$  times to each process!

# Reliable multicast over IP multicast

- IP multicast, piggybacked acknowledgements (in other messages) and negative acknowledgements
- Processes do not send separate acknowledgement messages
- Processes send a separate response message only when they detect that they have missed a message (negative ack)

# Reliable multicast over IP multicast

- Each process keeps
  - a sequence number (initially 0) for  $g$
  - ack per process: the latest message sequence number sent by others
- *R-multicast( $g, m$ )*
  - IP-multicast message to  $g$
  - include sequence number and acks from all processes
- Piggybacked acks enable recipients to learn about messages that they have not received



# Reliable multicast over IP multicast

- A process *R-delivers* a message destined for *g* carrying the sequence number *S* from *p* if and only if the stored ack for that process refers to  $S - 1$  (ack is incremented)
- Already delivered messages are discarded ( $ack > S$ )
- The message is stored in a hold-back queue
  - if the stored ack for that process refers to a value  $S' < S - 1$
  - or an acknowledgement in the message refers to a message from some other process that is not yet received

# Reliable multicast over IP multicast

- The process then sends a negative ack requesting the lost messages (there can be duplicates)
- When gaps are filled, it *R-delivers(m)* (removing it from queue)
- Assumption for agreement: messages flow indefinitely!!

# Uniform agreement

- Agreement
  - if a correct process delivers message  $m$ , then all correct processes in  $g$  will eventually deliver  $m$
- Uniform agreement
  - if a process, whether it is correct or fails, delivers message  $m$ , then all correct processes in  $g$  will eventually deliver  $m$
- R-multicast meets the uniform agreement condition
- Uniform agreement is important to keep state consistent! (e.g., replication)

# Ordered Multicast

- Some applications require ordering guarantees in group communication
- FIFO ordering
  - If a correct process issues  $\text{multicast}(g, m)$  and then  $\text{multicast}(g, m')$ , then every correct process that delivers  $m'$  will deliver  $m$  before  $m'$
- Causal ordering
  - If  $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ , where  $\rightarrow$  is the *happened-before* relation induced by messages sent between members of  $g$ , then any correct process that delivers  $m'$  will deliver  $m$  before  $m'$


# Ordered Multicast

- Total ordering
  - If a correct process delivers message  $m$  before  $m'$ , then any other correct process that delivers  $m'$  will deliver  $m$  before  $m'$
- Causal ordering implies FIFO ordering!
  - Any two multicasts by the same process are related by *happened-before*
- FIFO and Causal orderings are partial orderings
  - Not all messages are sent by the same process and some multicasts are concurrent (not ordered by *happened-before*)

# Ordered Multicast

- Total ordering is not necessarily a FIFO or causal ordering
- Ordering the delivery of multicast messages can be expensive in terms of delivery latency and bandwidth consumption

# Ordered Multicast



- **Notice**

- the consistent ordering of totally ordered messages  $T_1$  and  $T_2$
  - the FIFO-related messages  $F_1$  and  $F_2$  and
  - the casually related messages  $C_1$  and  $C_3$  and
  - The otherwise arbitrary delivery ordering of messages
- **Total Ordering doesn't say anything about how messages are ordered**

# Distributed Bulletin Board

Bulletin board: <i>os.interesting</i>		
Item	From	Subject
23	A.Hanlon	Mach
24	G.Joseph	Microkernels
25	A.Hanlon	Re: Microkernels
26	T.L'Heureux	RPC performance
27	M.Walker	Re: Mach
end		

- Posts as they appear to one of the users
- Reliable multicast is required if every user is to receive every posting
- At minimum, FIFO ordering is desirable
  - every posting from a given user will be received in the same order



# Distributed Bulletin Board

Bulletin board: <i>os.interesting</i>		
Item	From	Subject
23	A.Hanlon	Mach
24	G.Joseph	Microkernels
25	A.Hanlon	Re: Microkernels
26	T.L'Heureux	RPC performance
27	M.Walker	Re: Mach
end		

- Messages 25 and 27 appear after 24 and 23, respectively (causal ordering)
- The item numbering is the same for all users (Total ordering) – then users can discuss about items  $X$

# FIFO Ordering

- Sequence numbers just like in R-multicast over IP-multicast

# Causal Ordering

- *happened-before* relationship **only** as established by multicast messages
- Each process maintains a *vector clock*
  - entries count the number of multicast messages that *happened-before* the next multicast message
- CO-multicast
  - increase the process entry by 1 and *B-multicast*( $g, \langle m, T \rangle$ )

# Causal Ordering

- On B-deliver: place message in a hold-back queue until
  - all previous messages from the same process have been delivered
  - all previous messages that the sending process has delivered, are delivered (causally preceded – look at vector timestamps)

# Total Ordering (Sequencer)

- A special process is the **sequencer** (can be elected)
- TO-multicast( $g, m$ ): a unique identifier  $i$  is attached and then ***B-multicast( $g$  and sequencer,  $\langle m, i \rangle$ )***
- The *sequencer( $g$ )* maintains a group-specific sequence number to assign increasing and consecutive sequence numbers to the messages that it *B-delivers*

# Total Ordering (Sequencer)

- It announces the sequence numbers by *B-multicasting* **order** messages to *g*
- A message will sit in a process' hold-back queue until it can be *TO-delivered* according to its sequence number

# Total Ordering (Collective Agreement)

- TO-multicast: *B-multicasts(g, m)*, *m* carries a unique identifier *i*
- On *B-deliver(m)*
  - each process *p* sends to the sender a proposed sequence number, which is the *max(lastProposed, lastAgreed) + 1*, include process ID to break ties.
  - provisionally assigns the proposed value to the message in a hold-back queue
- Sender *B-multicasts(g, agreed\_sequence\_number)* : *maximum proposed*

# Total Ordering (Collective Agreement)

- On *B-deliver(agreed\_sequence\_number)*
  - Each process assigns the agreed value and reorders the message in the queue if the value differs from the proposed one (updates the *lastAgreed* value)
  - If message is at the head of the queue, then TO-deliver(m)
- Not causally or FIFO ordered delivery



# Next Lecture ...

- ✓ Introduction
- ✓ HTTP, Caching, and CDNs
- ✓ Views
- ✓ Templates
- ✓ Forms
- ✓ Models
- ✓ Security
- ✓ Transactions
- ✓ Remote Procedure Call
- ✓ Web Services
- ✓ Time
- ✓ Elections and Group Communication
- **Coordination and Agreement**