

Coordination and Agreement

Web Applications and Services
Spring Term

Naercio Magaia



Contents

- Distributed Mutual Exclusion
- Consensus

Distributed Mutual Exclusion

- Distributed processes need to coordinate their activities
- Mutual exclusion is required, when processes share resources, to prevent interference and ensure consistency
- Shared resources may reside on a server, or a collection of peer processes must coordinate their accesses amongst themselves
- An asynchronous distributed system of N processes $p_i, i = 1, 2, \dots, N$ that do not share variables – communicate through messaging

Distributed Mutual Exclusion

- Application-level protocol for executing a critical section (CS)
 - `enter()` (may block), `resourceAccess()`, `exit()`
- Essential requirements for mutual exclusion
 1. *Safety*: at most one process may execute in the CS at a time
 2. *Liveness*: requests to enter and exit the CS eventually succeed
 3. *Happened-before ordering*: if one request to enter the CS happened-before another, then entry to the CS is granted in that order

Distributed Mutual Exclusion


- Requirement 2 (i.e., *liveness*) implies freedom from deadlock and starvation
- Requirement 3 is about fairness: if all requests are related by happened-before ordering, then it is not possible for a process to enter the CS more than once while other waits to enter
- Evaluation criteria
 - number of messages sent for *enter* and *exit* operation
 - synchronisation delay (p_1 exiting and p_2 entering)

Central Server Algorithm

- A process sends a request message to the server and waits for a permission token (i.e., the reply)
- The server replies immediately, if no other process has the token at the time of the request, or queues the request if token is given
- When a process exits the CS, it sends the token back to the server
- If the queue of waiting processes is not empty, the server chooses the oldest entry in the queue, removes it and replies to the corresponding process


Central Server Algorithm

- *safety* and *liveness* requirements are met
- *happened-before ordering* requirement isn't
- Entering the CS takes two messages
- Exiting the CS takes one message
- Synchronization delay is a round-trip
- Failures?
- How to (s)**elect** the server?



Ring-based Algorithm

- No Central Server
- N processes arranged in a logical ring
- Each process has a communication channel to the next process
- A token is passed from process to process in a single direction
- If no access to CS is required, token is passed to next process
- A process that needs access to the CS waits for the token



Ring-based Algorithm

- Upon exiting the CS, the token is passed to the next process
- *liveness* and *safety* are met but not the *happened-before ordering*
 - Processes exchange messages independently of the rotation of the token
- Bandwidth is constantly consumed
 - Processes send messages around the ring even when no process requires access
- Delay to enter
 - 0 to N messages

Ring-based Algorithm

- Delay to exit
 - Only 1 message
- Synchronization delay
 - 1 to N messages

Mutual Exclusion with Multicast and Logical Clocks

- Processes multicast a request message
 - enter only when all the other processes have replied
- Each process keeps a Lamport Virtual Clock
 - updated for internal events and when receiving messages
- Messages requesting entry are of the form $\langle T, p_i \rangle$
- Each process records its state of being outside the CS (*RELEASED*), wanting entry (*WANTED*) or being in the CS (*HELD*) in a variable

Mutual Exclusion with Multicast and Logical Clocks

On initialization

state := RELEASED;

To enter the section

state := WANTED;

Multicast *request* to all processes;

T := request's timestamp;

Wait until (number of replies received = ($N - 1$));

state := HELD;

Request processing deferred here

On receipt of a request $\langle T_i, p_i \rangle$ at p_j ($i \neq j$)

if (*state* = HELD or (*state* = WANTED and $(T, p_j) < (T_i, p_i)$))

then

queue *request* from p_i without replying;

else

reply immediately to p_i ;

end if

To exit the critical section

state := RELEASED;

reply to any queued requests;

Mutual Exclusion with Multicast and Logical Clocks

- *Safety*


- If it were possible for two processes to enter the CS at the same time, both of those processes would have to have replied to the other ($\langle T_i, p_i \rangle$ are causally ordered)

- *Liveness*

- when exiting the CS, a process replies to all pending requests


- *Happened-before ordering*

- Lamport Clocks



Mutual Exclusion with Multicast and Logical Clocks

- Delay to enter
 - $2*(N - 1)$ messages or N if native multicast is supported
- Delay to exit
 - up to $N - 1$
- Synchronization delay
 - 1 message (i.e., one process waits for only 1 message)



Consensus

- A collection of processes p_i ($i = 1, 2, \dots, N$) communicating by message passing
- Requirement: consensus can be reached even in the presence of faults
 - communication is reliable
 - processes may fail (crash or Byzantine (i.e., arbitrary) process failures)
 - up to some number f of the N processes are faulty

Consensus

- To reach consensus, every process p_i begins in the undecided state and proposes a single value v_i , drawn from a set D
- Processes communicate by exchanging values
- In the end, each process sets the value of a decision variable d_i and enters the decided state

Consensus

- The requirements of a consensus algorithm is that the following conditions should hold for its every execution
 - Termination
 - Eventually each correct process sets its decision variable
 - Agreement
 - The decision value of all correct processes is the same
 - Integrity
 - If all correct processes proposed the same value, then any correct process in the decided state has chosen that value

Consensus

- Consensus in the absence of failures is straightforward!
 - Each process reliably multicasts a proposed value and waits for all other values
 - Selects the decided value using an appropriate method (majority, min or max)

Consensus in a Synchronous System

- Up to f of the N processes may exhibit crash failures. Timeouts are used to detect.

Algorithm for process $p_i \in g$; algorithm proceeds in $f+1$ rounds

On initialization

$Values_i^1 := \{v_i\}$; $Values_i^0 = \{\}$;

In round r ($1 \leq r \leq f+1$)

$B\text{-multicast}(g, Values_i^r - Values_i^{r-1})$; // Send only values that have not been sent

$Values_i^{r+1} := Values_i^r$;

while (in round r)

{

On B-deliver(V_j) from some p_j

$Values_i^{r+1} := Values_i^{r+1} \cup V_j$;

}

After $(f+1)$ rounds

Assign $d_i = \text{minimum}(Values_i^{f+1})$;

Consensus in a Synchronous System

- All correct processes will have all proposed values to decide in the end
- If a process failed while multicasting its value either one or none of the processes will circulate this value
- Termination is obvious, i.e., the system is synchronous
- To check the correctness of the algorithm, we must show that each process arrives at the same set of values at the end of the final round

Consensus in a Synchronous System

- Agreement and integrity will then follow
- Assume that two processes differ in their final set of values
- Some correct process p_i possesses a value v that another correct process p_j doesn't
- Only possible explanation: some other process p_k sent v to p_i and then crashed

Consensus in a Synchronous System

- But why p_k possesses v but not p_i ? Another process crashed in the previous round
- We assumed f crash failures and we have $f + 1$ rounds – a contradiction

Impossibility in asynchronous systems

- The assumption above was that message exchanges take place in rounds
 - ...processes are entitled to timeout and assume that a faulty process has not sent them a message within the round, because the maximum delay has been exceeded
- No algorithm can guarantee to reach consensus in an asynchronous system, even with one process crash failure
 - The proof is beyond the scope of this module

Impossibility in asynchronous systems

- In an asynchronous system, processes can respond to messages at arbitrary times, so a crashed process is indistinguishable from a slow one
- Note the word guarantee!
- Masking faults
- Consensus using failure detectors
 - Ignore processes that are considered failed (even if they are not)
 - Large timeouts

Next Lecture ...

- ✓ Introduction
- ✓ HTTP, Caching, and CDNs
- ✓ Views
- ✓ Templates
- ✓ Forms
- ✓ Models
- ✓ Security
- ✓ Transactions
- ✓ Remote Procedure Call
- ✓ Web Services
- ✓ Time
- ✓ Elections and Group Communication
- **Coordination and Agreement**